

Introduction to OpenMP

Christian Terboven, Dirk Schmidl

IT Center, RWTH Aachen University
{terboven,schmidl}@itc.rwth-aachen.de

- **De-facto standard for Shared-Memory Parallelization.**
- **1997: OpenMP 1.0 for FORTRAN**
- **1998: OpenMP 1.0 for C and C++**
- **1999: OpenMP 1.1 for FORTRAN (errata)**
- **2000: OpenMP 2.0 for FORTRAN**
- **2002: OpenMP 2.0 for C and C++**
- **2005: OpenMP 2.5 now includes both programming languages.**
- **05/2008: OpenMP 3.0 release**
- **07/2011: OpenMP 3.1 release**
- **07/2013: OpenMP 4.0 release**
- **11/2015: OpenMP 4.5 release**




<http://www.OpenMP.org>

RWTH Aachen University is a member of the OpenMP Architecture Review Board (ARB) since 2006.

Multi-Core System Architecture

Moore's Law still holds!



The number of transistors on a chip is still doubling every 24 months ...

... but the clock speed is no longer increasing that fast!

Instead, we will see many more cores per chip!

Source: Herb Sutter

www.gotw.ca/publications/concurrency-ddj.htm


Example for a SMP system

■ Dual-socket Intel Woodcrest (dual-core) system

- Two cores per chip, 3.0 GHz
- Each chip has 4 MB of L2 cache on-chip, shared by both cores
- No off-chip cache
- Bus: Frontsidebus


■ SMP: Symmetric Multi Processor

- Memory access time is uniform on all cores
- Limited scalability



OpenMP Overview & Parallel Region

■ OpenMP: Shared-Memory Parallel Programming Model.

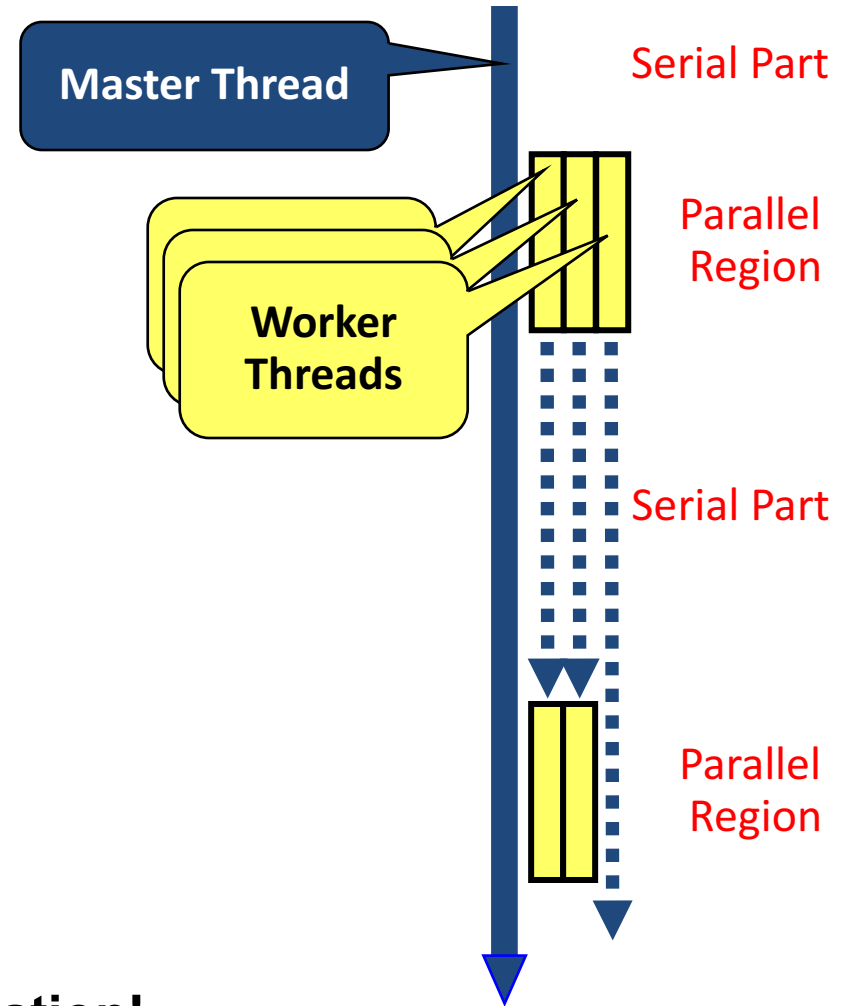


All processors/cores access a shared main memory.

Real architectures are more complex, as we will see later / as we have seen.

Parallelization in OpenMP employs multiple threads.

- OpenMP programs start with just one thread: The *Master*.
- *Worker* threads are spawned at *Parallel Regions*, together with the Master they form the *Team* of threads.
- In between Parallel Regions the Worker threads are put to sleep. The OpenMP *Runtime* takes care of all thread management work.
- Concept: *Fork-Join*.
- Allows for an incremental parallelization!



■ The parallelism has to be expressed explicitly.

C/C++

```
#pragma omp parallel
{
    ...
    structured block
    ...
}
```

Fortran

```
!$omp parallel
...
    structured block
...
$!omp end parallel
```

■ ***Structured Block***

- Exactly one entry point at the top
- Exactly one exit point at the bottom
- Branching in or out is not allowed
- Terminating the program is allowed
(abort / exit)

■ ***Specification of number of threads:***

- ▶ Environment variable:
OMP_NUM_THREADS=...
- ▶ Or: Via `num_threads` clause:
add `num_threads(num)` to the
parallel construct

Hello OpenMP World

Hello orphaned OpenMP World

- From within a shell, global setting of the number of threads:

```
export OMP_NUM_THREADS=4  
./program
```

- From within a shell, one-time setting of the number of threads:

```
OMP_NUM_THREADS=4 ./program
```

For Worksharing Construct

- If only the *parallel* construct is used, each thread executes the Structured Block.
- Program Speedup: *Worksharing*
- OpenMP's most common Worksharing construct: *for*

C/C++

```
int i;  
#pragma omp for  
for (i = 0; i < 100; i++)  
{  
    a[i] = b[i] + c[i];  
}
```

Fortran

```
INTEGER :: i  
!$omp do  
DO i = 0, 99  
    a[i] = b[i] + c[i];  
END DO
```

- Distribution of loop iterations over all threads in a Team.
- Scheduling of the distribution can be influenced.

- Loops often account for most of a program's runtime!

Pseudo-Code
Here: 4 Threads

Serial

```
do i = 0, 99  
  a(i) = b(i) + c(i)  
end do
```

Thread 1

```
do i = 0, 24  
  a(i) = b(i) + c(i)  
end do
```

Thread 2

```
do i = 25, 49  
  a(i) = b(i) + c(i)  
end do
```


Thread 3

```
do i = 50, 74  
  a(i) = b(i) + c(i)  
end do
```

Thread 4

```
do i = 75, 99  
  a(i) = b(i) + c(i)  
end do
```

Memory



Vector Addition

- **for-construct: OpenMP allows to influence how the iterations are scheduled among the threads of the team, via the *schedule* clause:**
 - ➔ `schedule(static [, chunk])`: Iteration space divided into blocks of chunk size, blocks are assigned to threads in a round-robin fashion. If chunk is not specified: #threads blocks.
 - ➔ `schedule(dynamic [, chunk])`: Iteration space divided into blocks of chunk (not specified: 1) size, blocks are scheduled to threads in the order in which threads finish previous blocks.
 - ➔ `schedule(guided [, chunk])`: Similar to dynamic, but block size starts with implementation-defined value, then is decreased exponentially down to chunk.
- **Default on most implementations is `schedule(static)`.**

■ Can all loops be parallelized with `for`-constructs? No!

→ Simple test: If the results differ when the code is executed backwards, the loop iterations are not independent. BUT: This test alone is not sufficient:

C/C++

```
int i, int s = 0;
#pragma omp parallel for
for (i = 0; i < 100; i++)
{
    s = s + a[i];
}
```

■ **Data Race:** If between two synchronization points at least one thread writes to a memory location from which at least one other thread reads, the result is not deterministic (race condition).

- **A *Critical Region* is executed by all threads, but by only one thread simultaneously (*Mutual Exclusion*).**

C/C++

```
#pragma omp critical (name)
{
    ... structured block ...
}
```

- **Do you think this solution scales well?**

C/C++

```
int i, s = 0;
#pragma omp parallel for
for (i = 0; i < 100; i++)
{
    #pragma omp critical
    { s = s + a[i]; }
}
```

Data Scoping

- **Managing the Data Environment is the challenge of OpenMP.**
- **Scoping in OpenMP: Dividing variables in *shared* and *private*:**
 - *private*-list and *shared*-list on Parallel Region
 - *private*-list and *shared*-list on Worksharing constructs
 - General default is *shared* for Parallel Region, *firstprivate* for Tasks.
 - Loop control variables on *for*-constructs are *private*
 - Non-static variables local to Parallel Regions are *private*
 - *private*: A new uninitialized instance is created for each thread
 - *firstprivate*: Initialization with Master's value
 - *lastprivate*: Value of last loop iteration is written back to Master
 - Static variables are *shared*

■ Global / static variables can be privatized with the *threadprivate* directive

- One instance is created for each thread
 - Before the first parallel region is encountered
 - Instance exists until the program ends
 - Does not work (well) with nested Parallel Region
- Based on thread-local storage (TLS)
 - TlsAlloc (Win32-Threads), pthread_key_create (Posix-Threads), keyword `__thread` (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

Fortran

```
SAVE INTEGER :: i  
!$omp threadprivate(i)
```

■ Global / static variables can be privatized with the *threadprivate* directive

- One instance is created for each thread
 - Before the first parallel region is encountered
 - Instance exists until the program ends
 - Does not work (well) with nested Parallel Region
- Based on thread-local storage (TLS)
 - TlsAlloc (Win32-Threads), pthread_key_create (Posix-Threads), keyword `__thread` (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

Fortran

```
SAVE INTEGER :: i  
!$omp threadprivate(i)
```

The Barrier Construct

- **OpenMP `barrier` (implicit or explicit)**

→ Threads wait until all threads of the current *Team* have reached the barrier

C/C++

```
#pragma omp barrier
```

- **All worksharing constructs contain an implicit barrier at the end**

Back to our bad scaling example

C/C++

```
int i, s = 0;
#pragma omp parallel for
for (i = 0; i < 100; i++)
{
    #pragma omp critical
    { s = s + a[i]; }
}
```

It's your turn: Make It Scale!



```
#pragma omp parallel
```

```
{
```

```
#pragma omp for
```

```
  for (i = 0; i < 99; i++)
```

```
  {
```

```
      s = s + a[i];
```

```
  }
```

```
} // end parallel
```

```
do i = 0, 99  
  s = s + a(i)  
end do
```



```
do i = 0, 24  
  s = s + a(i)  
end do
```

```
do i = 25, 49  
  s = s + a(i)  
end do
```

```
do i = 50, 74  
  s = s + a(i)  
end do
```

```
do i = 75, 99  
  s = s + a(i)  
end do
```

- In a *reduction*-operation the operator is applied to all variables in the list. The variables have to be *shared*.

→ `reduction(operator:list)`

→ The result is provided in the associated reduction variable

C/C++

```
int i, s = 0;

#pragma omp parallel for reduction(+:s)
for(i = 0; i < 99; i++)
{
    s = s + a[i];
}
```


→ Possible reduction operators with initialization value:

`+` (0), `*` (1), `-` (0), `&` (~0), `|` (0), `&&` (1), `||` (0),


`^` (0), `min` (largest number), `max` (least number)

```
double s_priv[nthreads];  
  
#pragma omp parallel num_threads(nthreads)  
{  
    int t=omp_get_thread_num();  
  
    #pragma omp for  
    for (i = 0; i < 99; i++)  
    {  
        s_priv[t] += a[i];  
    }  
} // end parallel  
for (i = 0; i < nthreads; i++)  
{  
    s += s_priv[i];  
}
```

- When data is used, it is copied into caches.
- The hardware always copies chunks into the cache, so called *cache-lines*.
- This is useful, when:
 - the data is used frequently (temporal locality)
 - consecutive data is used which is on the same cache-line (spatial locality)




- **False Sharing occurs when**
 - different threads use elements of the same cache-line
 - one of the threads writes to the cache-line
- **As a result the cache line is moved between the threads, also there is no real dependency**
- **Note: False Sharing is a performance problem, not a correctness issue**




False Sharing

- no performance benefit for more threads
- Reason: false sharing of `s_priv`
- Solution: padding so that only one variable per cache line is used



— with false sharing — without false sharing




```
double s_priv[nthreads * 8];

#pragma omp parallel num_threads(nthreads)
{
    int t=omp_get_thread_num();

    #pragma omp for
    for (i = 0; i < 99; i++)
    {
        s_priv[t * 8] += a[i];
    }
} // end parallel
for (i = 0; i < nthreads; i++)
{
    s += s_priv[i * 8];
}
```

PI

Example: Pi (1/2)




```
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}
```

```
double CalcPi (int n)
{
    const double fH  = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;

    #pragma omp parallel for
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}
```

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



Example: Pi (1/2)




```
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}
```

```
double CalcPi (int n)
{
    const double fH  = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;

    #pragma omp parallel for private(fX,i) reduction(+:fSum)
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}
```

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



■ Results:

# Threads	Runtime [sec.]	Speedup
1	1.11	1.00
2		
4		
8	0.14	7.93

■ Scalability is pretty good:

- About 100% of the runtime has been parallelized.
- As there is just one parallel region, there is virtually no overhead introduced by the parallelization.
- Problem is parallelizable in a trivial fashion ...

Single and Master Construct

C/C++

```
#pragma omp single [clause]  
... structured block ...
```

Fortran

```
!$omp single [clause]  
... structured block ...  
!$omp end single
```

- The `single` construct specifies that the enclosed structured block is executed by only one thread of the team.

→ It is up to the runtime which thread that is.

- Useful for:

→ I/O

→ Memory allocation and deallocation, etc. (in general: setup work)

→ Implementation of the single-creator parallel-executor pattern as we will see now...

C/C++

```
#pragma omp master[clause]  
... structured block ...
```

Fortran

```
!$omp master[clause]  
... structured block ...  
!$omp end master
```

- The `master` construct specifies that the enclosed structured block is executed only by the master thread of a team.
- Note: The master construct is no worksharing construct and does not contain an implicit barrier at the end.

Section and Ordered Construct

■ How would you parallelize this code?

```
void traverse (Tree *tree)
{
    if (tree->left)    traverse(tree->left) ;

    if (tree->right)    traverse(tree->right) ;

    process (tree) ;
}
```

■ One option: Use OpenMP's parallel sections.

C/C++

```
#pragma omp sections [clause]
{
    #pragma omp section
    ... structured block ...
    #pragma omp section
    ... structured block ...
    ...
}
```

Fortran

```
!$omp sections [clause]
!$omp section
... structured block ...
!$ omp section
... structured block ...
...
!$omp end sections
```

- The **sections** construct contains a set of structured blocks that are to be distributed among and executed by the team of threads.

■ How would you parallelize this code?

```
void traverse (Tree *tree)
{
#pragma omp parallel sections
{
#pragma omp section
    if (tree->left)    traverse(tree->left) ;
#pragma omp section
    if (tree->right)   traverse(tree->right) ;
} // end omp parallel
    process(tree) ;
}
```

Nested Parallel Regions

Barrier here!

■ Downsides of this option:

- Unnecessary overhead and synchronization points
- Not always well supported (how many threads to be used?)

- **Allows to execute a structured block within a parallel loop in sequential order**

→ In addition, an `ordered` clause has to be added to the *for* construct which any *ordered* construct may occur

```
#pragma omp parallel for ordered
for (i=0 ; i<10 ; i++){
    ...
    #pragma omp ordered
    {
        ...
    }
    ...
}
```

- **Use Cases:**

→ Can be used e.g. to enforce ordering on printing of data

→ May help to determine whether there is a data race

Runtime Library

■ C and C++:

- If OpenMP is enabled during compilation, the preprocessor symbol `_OPENMP` is defined. To use the OpenMP runtime library, the header `omp.h` has to be included.
- `omp_set_num_threads(int)`: The specified number of threads will be used for the parallel region encountered next.
- `int omp_get_num_threads`: Returns the number of threads in the current team.
- `int omp_get_thread_num()`: Returns the number of the calling thread in the team, the Master has always the id 0.

■ Additional functions are available, e.g. to provide locking functionality.

Tasking

Recursive approach to compute Fibonacci



```
int main(int argc,  
        char* argv[])  
{  
    [...]  
    fib(input);  
    [...]  
}
```

```
int fib(int n)    {  
    if (n < 2) return n;  
    int x = fib(n - 1);  
    int y = fib(n - 2);  
    return x+y;  
}
```

- On the following slides we will discuss three approaches to parallelize this recursive code with Tasking.

C/C++

```
#pragma omp task [clause]  
... structured block ...
```

Fortran

```
!$omp task [clause]  
... structured block ...  
!$omp end task
```

■ Each encountering thread/task creates a new Task

- Code and data is being packaged up
- Tasks can be nested
 - Into another Task directive
 - Into a Worksharing construct

■ Data scoping clauses:

- `shared(list)`
- `private(list)` `firstprivate(list)`
- `default(shared | none)`

- **Some rules from *Parallel Regions* apply:**
 - Static and Global variables are shared
 - Automatic Storage (local) variables are private

- **If shared scoping is not derived by default:**
 - Orphaned Task variables are `firstprivate` by default!
 - Non-Orphaned Task variables inherit the `shared` attribute!
 - Variables are `firstprivate` unless `shared` in the enclosing context

First version parallelized with Tasking (omp-v1)




```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```


```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- **Only one Task / Thread enters `fib()` from `main()`, it is responsible for creating the two initial work tasks**
- **Taskwait is required, as otherwise `x` and `y` would be lost**


- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks



Task Queue




- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks
- ...



■ Overhead of task creation prevents scalability!

Speedup of Fibonacci with Tasks



- If the expression of an **if** clause on a task evaluates to **false**
 - The encountering task is suspended
 - The new task is executed immediately
 - The parent task resumes when the new task finishes
 - Used for optimization, avoids queuing of small tasks


- **Improvement: Don't create yet another task once a certain (small enough) n is reached**

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
        [...]
    }
```

```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x) \
        if(n > 30)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y) \
        if(n > 30)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

■ Speedup is better, but still not great

Speedup of Fibonacci with Tasks




■ Improvement: Skip the OpenMP overhead once a certain *n* is reached

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    if (n <= 30)
        return serfib(n);
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```


■ Looks promising...

Speedup of Fibonacci with Tasks




- First two versions were slow because of overhead!

Runtime of Fibonacci with Tasks



- Third version is comparable to serial version w/o OpenMP 😊

Runtime of Fibonacci with Tasks



Data Scoping Example (1/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a:
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```

Data Scoping Example (2/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```


Data Scoping Example (3/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c:
            // Scope of d:
            // Scope of e:
        }
    }
}
```

Data Scoping Example (4/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d:
            // Scope of e:
        }
    }
}
```

Data Scoping Example (5/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e:

        }
    }
}
```

Data Scoping Example (6/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e: private
        }
    }
}
```

Hint: Use default(none) to be forced to think about every variable if you do not see clear.

Data Scoping Example (7/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared,          value of a: 1
            // Scope of b: firstprivate,    value of b: 0 / undefined
            // Scope of c: shared,          value of c: 3
            // Scope of d: firstprivate,    value of d: 4
            // Scope of e: private,         value of e: 5
        }
    }
}
```

■ OpenMP `barrier` (implicit or explicit)

- All tasks created by any thread of the current *Team* are guaranteed to be completed at barrier exit

C/C++

```
#pragma omp barrier
```

■ Task barrier: `taskwait`

- Encountering Task suspends until child tasks are complete
 - Only direct childs, not descendants!

C/C++

```
#pragma omp taskwait
```

■ Task Synchronization explained:

```
#pragma omp parallel num_threads(np)
{
    #pragma omp task
        function_A();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
            function_B();
    }
}
```

np Tasks created here, one for each thread

All Tasks guaranteed to be completed here

1 Task created here

B-Task guaranteed to be completed here

Questions?